



Trinity College Dublin

Coláiste na Tríonóide, Baile Átha Cliath

The University of Dublin

CSU44000 Internet Applications

Week 6 Lecture 1

Conor Sheedy

Practical Cloud Computing with Amazon Web Services

Amazon Web Services, AWS

Offer a wide range of services

Original product was Infrastructure-as-a-service

- **EC2**
- **Many services built on top of that**

Amazon Elastic Compute Cloud

Virtual Machine Types

- Hardware
- Software

Availability Zones & Regions

Pricing

- Competitiveness with running own infrastructure

Security Groups

Attached Storage : Elastic Block Store

The Virtual Machine Instance

User chooses to instantiate a virtual machine of a given type of hardware with a given type of software

- **Depends on application needs**
 - Eg. Fast CPU?
 - Lot of memory?

<https://aws.amazon.com/ec2/instance-types/>

- **CPUs can ‘burst’ using a system of CPU credits/hour**
 - instances accumulate credits when idle and consume them when busy
- **Instances are allocated 3,6,12... CPU credits/hour**

We will mostly use:

Instance	vCPU*	CPU Credits / hour	Mem (GiB)	Storage	Network Performance
t2.nano	1	3	0.5	EBS-Only	Low

- **Each Virtual CPU (vCPU) is equivalent of one Intel Xeon Core or Hyperthread (depending on instance)**
- **Memory/Networking depends on instance type**

The Machine Image

- Each machine instance is created with a local enduring storage
- The operating system and installed software constitute the Amazon Machine Image (AMI)
- Software licence fees can be included in the cost of running an instance
- ‘Curated’ software stacks are available by ‘subscribing’ to sellers in the AMI Marketplace (<https://aws.amazon.com/marketplace>)
- Users can create their own machine image with the software mix they need

Sample Pricing – EU(Ireland) 2019

vCPU	ECU	Memory (GiB)	Instance Storage (GB)	Linux/UNIX Usage	
t3.nano	2	Variable	0.5 GiB	EBS Only	\$0.0057 per Hour
t3.micro	2	Variable	1 GiB	EBS Only	\$0.0114 per Hour
t3.small	2	Variable	2 GiB	EBS Only	\$0.0228 per Hour
t3.medium	2	Variable	4 GiB	EBS Only	\$0.0456 per Hour
p3dn.24xlarge	96	345	768 GiB	2 x 900 NVMe SSD	\$33.711 per Hour
i3en.24xlarge	96	N/A	768 GiB	8 x 7500 NVMe SSD	\$12.00 per Hour

Cost Competitiveness

- **Can be a low cost option for businesses compared with running their own hardware**
- **Cost control becomes an important concern**
- **Ease of scaling could lead to costs scaling**
- **Can be a large fraction of total cost for some businesses**

Disk Storage

- **Local hard disk(or SSD) is provided by Elastic Block Store (EBS) images**
- **By default each created machine instance gets one (LINUX AMI defaults to 8 GiB)**
- **After an instance is stopped, user is still charged rental for the EBS store – until the instance is terminated**
- **Charges are based on location and space used over time**
- **E.g. 2019, Europe(Ireland)**

\$0.11 per GB-month of provisioned storage

Launching the Instance – Security Environment

- Each machine instance is launched within a ‘Security Group’
- By default everything is closed – unless explicitly opened
- Security group can be created ‘live’ using the launch wizard or pre-configured

Step 6: Configure Security Group

A security group is a set of firewall rules that control the traffic for your instance. On this page, you can add rules to allow specific traffic to reach your instance. For example, if you want to set up a web server and allow Internet traffic to reach your instance, add rules that allow unrestricted access to the HTTP and HTTPS ports. You can create a new security group or select from an existing one below. [Learn more](#) about Amazon EC2 security groups.

Assign a security group: ☒ Create a new security group

☐ Select an existing security group

Security group name:

Description:

Type <small>i</small>	Protocol <small>i</small>	Port Range <small>i</small>	Source <small>i</small>	Description <small>i</small>	
SSH ▾	TCP	22	Custom ▾ 0.0.0.0/0	e.g. SSH for Admin Desktop	✕
HTTP ▾	TCP	80	Custom ▾ 0.0.0.0/0	e.g. SSH for Admin Desktop	✕
Custom TCP F ▾	TCP	3000	Anywhere ▾ 0.0.0.0/0, ::/0	Allow 3000 for testing...	✕
<div>Add Rule</div>					

Cancel

Previous

Review and Launch

AWS Student Environment

- **AWS recently updated learning support tools**
 - <https://aws.amazon.com/education/awseducate/>
 - Register now with tcd email
- **Sandboxed environment**
 - No such thing a 'free'
 - Free Tier runs out
 - Credit card not needed in a Sandbox
 - Otherwise Credit card required
 - Mistakes can be costly
 - On gaining sufficient experience in the sandbox
 - Sign up at your own risk
 - Manage that risk appropriately

Storage Hierarchy

Elastic Block Store (EBS)

- Fast storage
- Can specify speed of access
- Indented for short term storage

Simple Storage Service (S3)

- Standard for longer term, larger storage
- Various flavours based on frequency or speed of Access

Glacier – a slow cheap (tape-based) storage for end-of-life data (perhaps held for compliance purposes)

- Suitable for archival storage

Simple Storage Service, S3

- **Introduced in 2006**
- **Used to store Objects – the Objects are grouped into Buckets**
- **Much Cheaper than using Elastic Block Storage, EBS:**

	Pricing
S3 Standard Storage	
First 50 TB / Month	\$0.023 per GB
Next 450 TB / Month	\$0.022 per GB
Over 500 TB / Month	\$0.021 per GB

- **Buckets are given names**
 - must be unique across AWS

S3 Buckets Used by

- **Dropbox up to 2016**
 - Freemium business model
 - Loss made on users of free
 - Needs to be covered by paying customers
- **Netflix as ‘System of Record’**
- **Tumblr, Pinterest**
- **Can be configured to serve static web file**
 - cheap and effective static website hosting

S3 Bucket Contents

- **A bucket can contain objects identified by a 'Key' = name**
- **The bucket has no internal structure (bag of objects), but the AWS interface supports folders by using '/' in the name**
- **If you upload a cat_photo.jpg into the folder /myphotos**
 - it uploads and creates an object with the name /myphotos/cat_photo.jpg
- **Access Control**
 - Can be applied at Bucket and at Object Level
 - Can be accessed using a REST style CRUD interface using: <http://bucket.s3-aws-region.amazonaws.com> or s3-aws-region.amazonaws.com/bucket
 - Wrapped API available (AWS-SDK) for Node.js

S3 Usage Pricing

PUT, COPY, POST, or LIST Requests	\$0.005 per 1,000 requests
GET, SELECT and all other Requests	\$0.0004 per 1,000 requests
Lifecycle Transition Requests into Standard - Infrequent Access or One Zone - Infrequent Access or Intelligent-Tiering	\$0.01 per 1,000 requests

Node.js code to upload a file to an S3 bucket

```
/ Load the AWS SDK for Node.js
var AWS = require('aws-sdk');
// Set the region
AWS.config.update({region: 'REGION'});
```

```
// Create S3 service object
s3 = new AWS.S3({apiVersion: '2006-03-01'});
```

```
// call S3 to retrieve upload file to specified bucket
// process gets parameters from the command line passed
// to node.js
var uploadParams = {Bucket: process.argv[2], Key: ' ', Body: '
'};
var file = process.argv[3];
```

```
// Configure the file stream and obtain the upload parameters
var fs = require('fs');
var fileStream = fs.createReadStream(file);
fileStream.on('error', function(err) {
  console.log('File Error', err);
});
uploadParams.Body = fileStream;
var path = require('path');
uploadParams.Key = path.basename(file);
```

```
// call S3 to retrieve upload file to specified bucket
s3.upload(uploadParams, function (err, data) {
  if (err) {
    console.log("Error", err);
  } if (data) {
    console.log("Upload Success", data.Location);
  }
});
```

node s3_upload.js BUCKETNAME Filename

Archive Storage - Glacier

A slow cheap (tape-based) storage for end-of-life data (perhaps held for compliance purposes)



Glacier Storage Costs	\$0.004 per GB/month
Retrieval Pricing (expedited)	0.01 (0.03) per GB
Retrieval Request Pricing	0.055 per 1,000 requests

Geographic Distribution


- **Amazon divides the world into 16 Geographic Regions**
 - e.g. US East, EU(Frankfurt)
- **Each region contains at least two availability zones**
 - These are essentially data center locations
 - If low latency is required, Region should be close to user
- **Things (servers, storage) are often replicated across availability zones**
- **Regions are distinct unless the user explicitly connects them**

Geographic Distribution

Monthly cost:

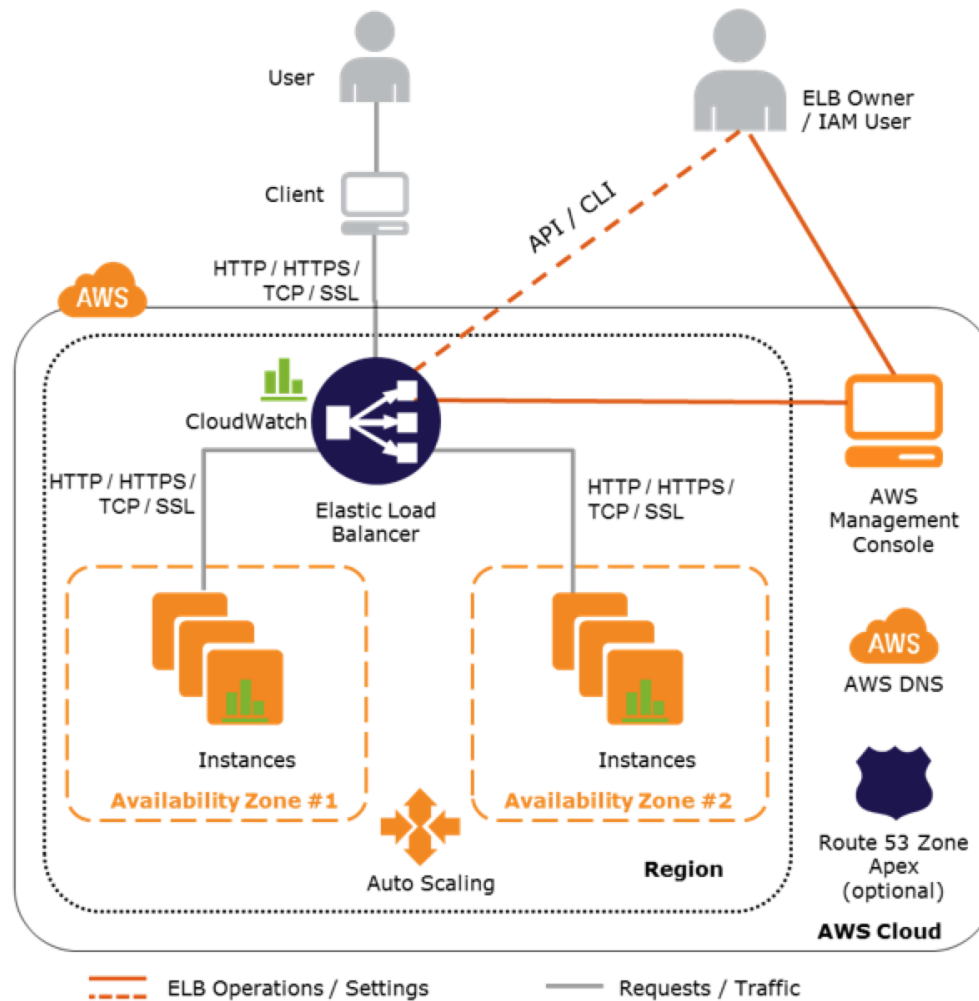
- **Pricing often varies a lot between regions**
 - Spot pricing
 - Unused capacity
 - <https://aws.amazon.com/ec2/spot/>
- **Factors in Choosing a location**
 - If low latency is required
 - Region should be close to user

Cost by AWS region: 5 c5.large, 20GB gp2 EBS storage each, 1 ELB, 5.1TB data processed



Sharing load across Instances

- **Machine Instances can be added to a load-balancing group which means their load is passed through a LB node**
- **Application level load balancing**
 - Done at OSI layer 7
 - each HTTP request is routed to different nodes
- **Network Level load Balancing**
 - Incoming TCP connections/UDP are routed to different nodes
- **AWS Video on elastic load balancer**
https://www.youtube.com/watch?v=YO4L_9poF3g



Scaling a Service

- **Manually**
- **Establish an Auto-Scaler and add VMs to the group**
- **Establish Launch Configuration**
 - Use an AMI ‘Golden Image’
- **Monitoring – define what constitutes a ‘Health Check’**
 - EC2 Instance Status
 - Fetching a particular URL
- **Autoscaler will terminate unhealthy instances**
- **Will start more to maintain ‘desired’ number of instances**
- **Common to use multiple ‘availability zones’**
 - so your service will survive failure of an entire data center

Cloud Database Services

- **Database Speed, Reliability and Cost are major pillars of most networked applications**
- **Historically people have used their own privately run databases based on technology like SQL or noSQL**
- **In migrating to the cloud, this is a major choice**
- **Difficult to change once the choice is made**
 - Vendor Lock-in
 - Migration issues are a significant consideration
- **Each major cloud provider prefers a different technology**
 - Amazon AWS DynamoDB (also have AWS RDS for relational)
 - Google Cloud BigTable
 - Microsoft SQL

AWS Database Services - DynamoDB

- **noSQL Database service**
- **At highest level, DynamoDB is a Key:Value Store**
- **It stores items in tables with one or a combination of attributes stored as a primary key**
- **Tables are partitioned and replicated for high availability, reliability and performance**
- **Architecture described in 2007 SOSP paper:Dynamo: Amazon's Highly Available Key-value Store**
<https://www.allthingsdistributed.com/files/amazon-dynamo-sosp2007.pdf>
- **Uses Leslie Lamport's Paxos algorithm to achieve consensus among partitions**
<https://lamport.azurewebsites.net/pubs/paxos-simple.pdf>

AWS DynamoDB Tutorial

https://www.youtube.com/watch?v=2mVR_Qgx_RU (9 minutes)

DynamoDB Points to Note

- **It's a NoSQL database with very high performance**
 - can handle 10 trillion requests per day and support peaks of more than 20 million requests per second
- **Read Performance is very high (and cheap)**
- **Write Performance lower (and more expensive)**
- **Highly Replicated with “Eventual Consistency” by default**
 - Can switch on “Strong Consistency” but lower performance and more expensive

Accessing Dynamodb from Javascript

- **Amazon provides a wrapped API `aws-sdk` to access the DynamoDB service**

```

// Load the AWS SDK for Node.js
var AWS = require('aws-sdk');
// Set the region
AWS.config.update({region: 'REGION'});

// Create the DynamoDB service object
var ddb = new AWS.DynamoDB({apiVersion: '2012-08-10'});

var params = {
  AttributeDefinitions: [
    {
      AttributeName: 'CUSTOMER_ID',
      AttributeType: 'N'
    },
    {
      AttributeName: 'CUSTOMER_NAME',
      AttributeType: 'S'
    }
  ],

```

```

  KeySchema: [
    {
      AttributeName: 'CUSTOMER_ID',
      KeyType: 'HASH'
    },
    {
      AttributeName: 'CUSTOMER_NAME',
      KeyType: 'RANGE'
    }
  ],
  ProvisionedThroughput: {
    ReadCapacityUnits: 1,
    WriteCapacityUnits: 1
  },
  TableName: 'CUSTOMER_LIST',
  StreamSpecification: {
    StreamEnabled: false
  }
};

// Call DynamoDB to create the table
ddb.createTable(params, function(err, data) {
  if (err) {
    console.log("Error", err);
  } else {
    console.log("Table Created", data);
  }
});

```

Exercise

- **Create a DynamoDB Database**
- **Populate it with items taken from an S3 bucket**
- **Allow Simple queries**
- **Destroy the database**

```
KeySchema: [  
  {  
    AttributeName: 'CUSTOMER_ID',  
    KeyType: 'HASH'  
  },  
  {  
    AttributeName: 'CUSTOMER_NAME',  
    KeyType: 'RANGE'  
  }  
],  
ProvisionedThroughput: {  
  ReadCapacityUnits: 1,  
  WriteCapacityUnits: 1  
},  
TableName: 'CUSTOMER_LIST',  
StreamSpecification: {  
  StreamEnabled: false  
}  
};  
  
// Call DynamoDB to create the table  
ddb.createTable(params, function(err, data) {  
  if (err) {  
    console.log("Error", err);  
  } else {  
    console.log("Table Created", data);  
  }  
});
```