



**Trinity College Dublin**

Coláiste na Tríonóide, Baile Átha Cliath

The University of Dublin

# CSU44000 Internet Applications

Week 10 Lecture 1, Serverless

Conor Sheedy

# Container-as-a-Service

- **Amazon Elastic Container Service (ECS).**
  - Your container instance must be running the Amazon ECS container agent.
  - The container agent is able to register the instance into one of your clusters.
  - If you use an Amazon ECS-optimized AMI the agent is already installed.
  - To use a different operating system you must install the agent.
  - An ECS container instance is an EC2 instance that runs the ECS Container Agent.
  - The EC2 instance is owned and managed by you.
  - The instance appears in the list of EC2 instances like any other EC2 instance.

# Container-as-a-Service

- **Amazon Elastic Container Service (ECS).**
  - Usually, you run a cluster of container instances in an auto-scaling group.
  - ECS is free of charge.
  - You only pay for the EC2 instances.
  - The downside is that you have to scale, monitor, patch, and secure the EC2 instances yourself.
  - An ECS container instance can run on Linux or Windows.
  - Unused CPU shares can be used by other containers if available.

# Serverless Computing – Container-as-a-Service

- **Fargate**
  - users are freed from thinking about the nature of machine resources that should be reserved for the task.
  - Instead, users just draw the resources they need from a seemingly infinite pool
  - AWS Fargate manages the task execution.
  - No EC2 instances to manage.
  - You pay for running tasks.
- **Serverless**
  - there are servers involved, but service users don't need to be concerned with them.
- **The packaging of functionality is a LINUX container**
  - still has quite a large granularity

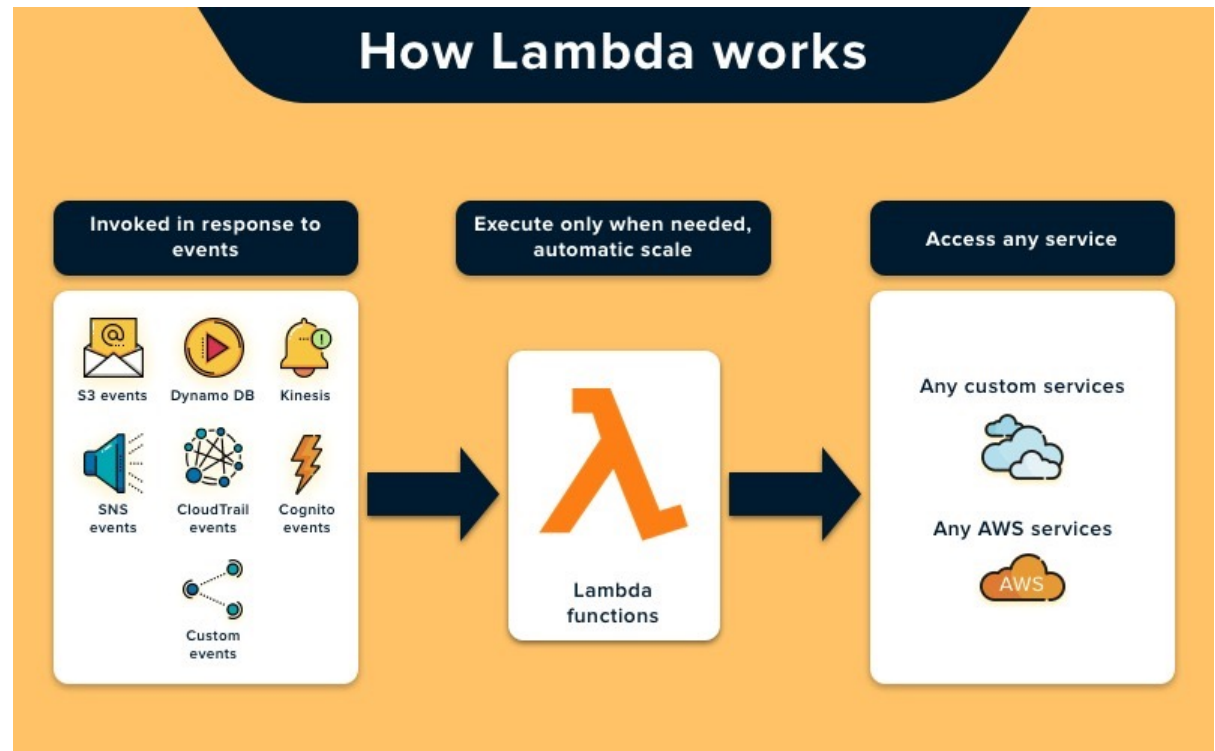
# Serverless Computing – Function-as-a-Service

- **AWS Lambda**
  - Event driven
  - Serverless
  - November 2014
  - Node.js, Python, Java, Go, Ruby, C#
  - Name comes from “Lambda functions”, anonymous functions
- **Google Cloud Functions**
- **Azure Functions**

**Reduces complexity to the bare minimum – a single function!**

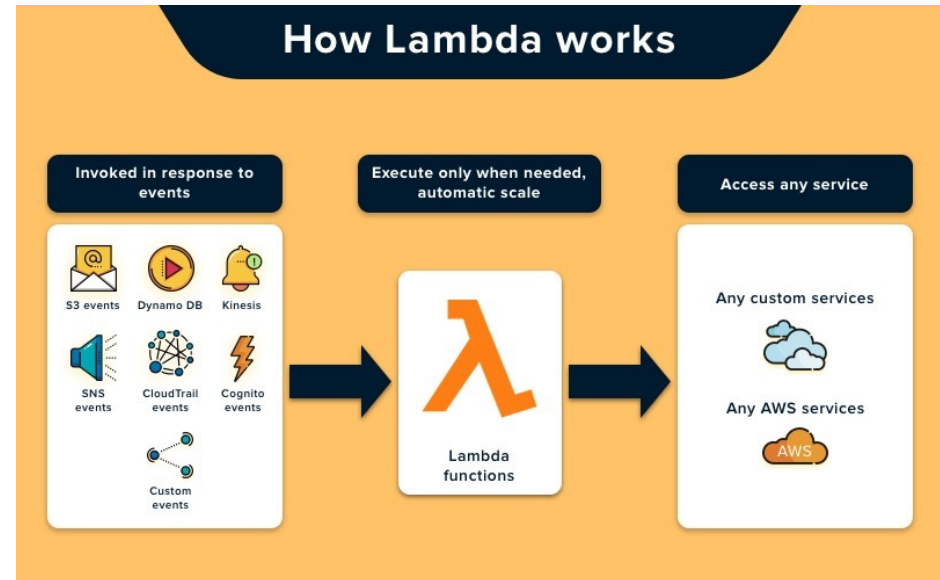
# The AWS Lambda Architecture

- A variety of services - including incoming http requests trigger AWS Lambda requests
- Code is launched into a container on an EC2 instance with an Amazon Linux AMI
  - The point is that the implementation is not the developers concern
- Event is passed into the code



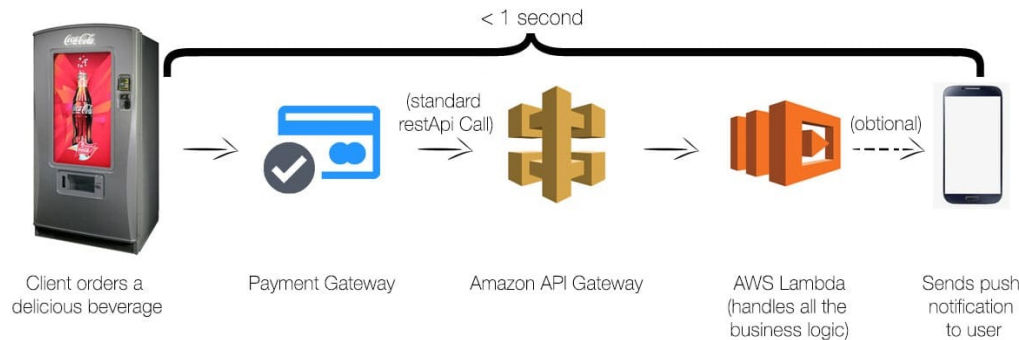
# The AWS Lambda Use Cases

- **Event Driven**
  - handling image and object uploads to Amazon S3
  - updating DynamoDB tables
  - responding to website clicks
  - reacting to sensor readings from IoT connected devices
- **custom HTTP requests**
  - configured in AWS API Gateway
  - which can handle authentication and authorisation



# Case Study: CocaCola's use of Serverless

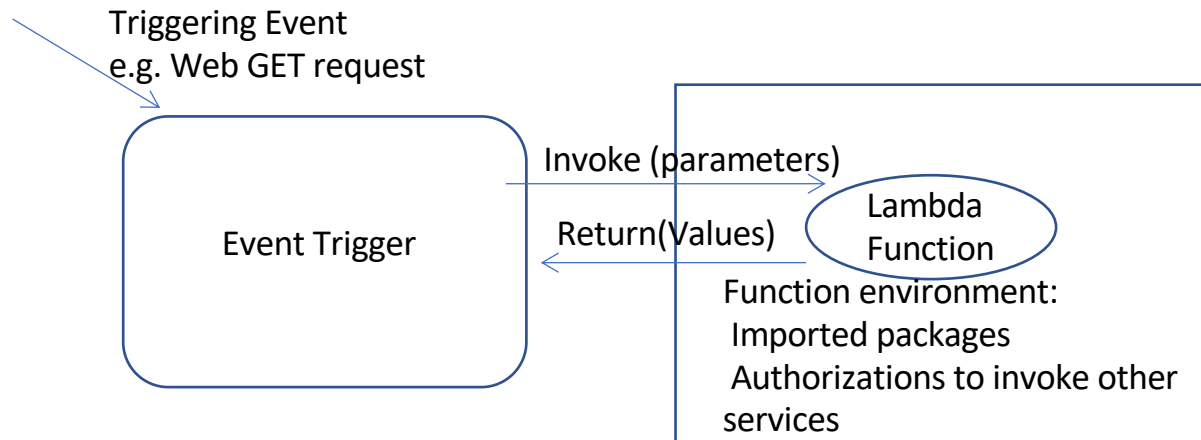
<https://dashbird.io/blog/serverless-case-study-coca-cola/>



- **Up to 2016 using 6 EC2 T2.medium**
  - cost \$12,864 per year
- **Serverless**
  - \$4,490 per year
  - 30 million requests per month
- **Infrastructure as a Service would start to get cheaper at 80 million requests per month**



# Lambda Function Components



- **Much of the complexity in Serverless is not in writing the Lambda Function, but in deploying it in the right environment, with correct event triggers and authorizations**
- **Frameworks like Claudia.js and “Serverless” simplify this**
  - create functioning Node. Js program, locally, then invoke the “Deploy” step which packages everything into a .ZIP file and uploads to AWS Lambda
- **We may use a more low-level approach for educational purposes, but hard to get anything complex done**
  - In practice use JavaScript Frameworks which hide the complexity
  - Reduce the time to develop and deploy

# Context

Events are passed in with a *Context* & *an Event*

We will look at an example event:

**GET /lambda?query=1234ABCD**

**Which arrives at an Application Load Balancer**

## Example Context & Event from an Application Load Balancer

```
{
  "requestContext": {
    "elb": {
      "targetGroupArn": "arn:aws:elasticloadbalancing:us-east-2:123456789012:targetgroup/lambda-279XGJDqGZ5rsrHC2Fjr/49e9d65c45c6791a"
    }
  },
  "httpMethod": "GET",
  "path": "/lambda",
  "queryStringParameters": {
    "query": "1234ABCD"
  },
  "headers": {
    "accept": "text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,image/apng,*/*;q=0.8",
    "accept-encoding": "gzip",
    "accept-language": "en-US,en;q=0.9",
    "connection": "keep-alive",
    "host": "lambda-alb-123578498.us-east-2.elb.amazonaws.com",
    "upgrade-insecure-requests": "1",
    "user-agent": "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/71.0.3578.98 Safari/537.36",
    "x-amzn-trace-id": "Root=1-5c536348-3d683b8b04734faae651f476",
    "x-forwarded-for": "72.12.164.125",
    "x-forwarded-port": "80",
    "x-forwarded-proto": "http",
    "x-imforwards": "20"
  },
  "body": "",
  "isBase64Encoded": false
}
```

## Exercise:

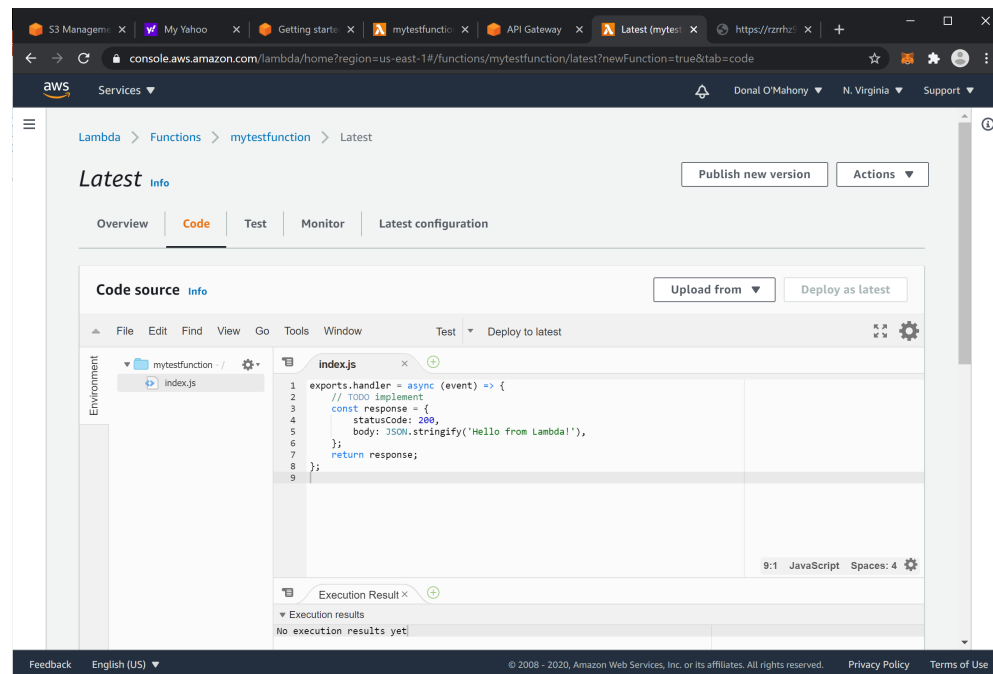
Create a simple function on AWS Lambda that sends the event parameters back in JSON format

- **The Lambda function is invoked and passed in an event**
- **If it is a HTTP GET, then it passes back JSON containing all the event parameters**

```
exports.handler = async (event) => {  
  let response = { statusCode: 200,  
    body: JSON.stringify('Hello from Lambda!'),  };  
  
  if (event.httpMethod == 'GET')  {  
    response = { statusCode :200,  
      body: JSON.stringify(event)}}  
  
  return response;}
```

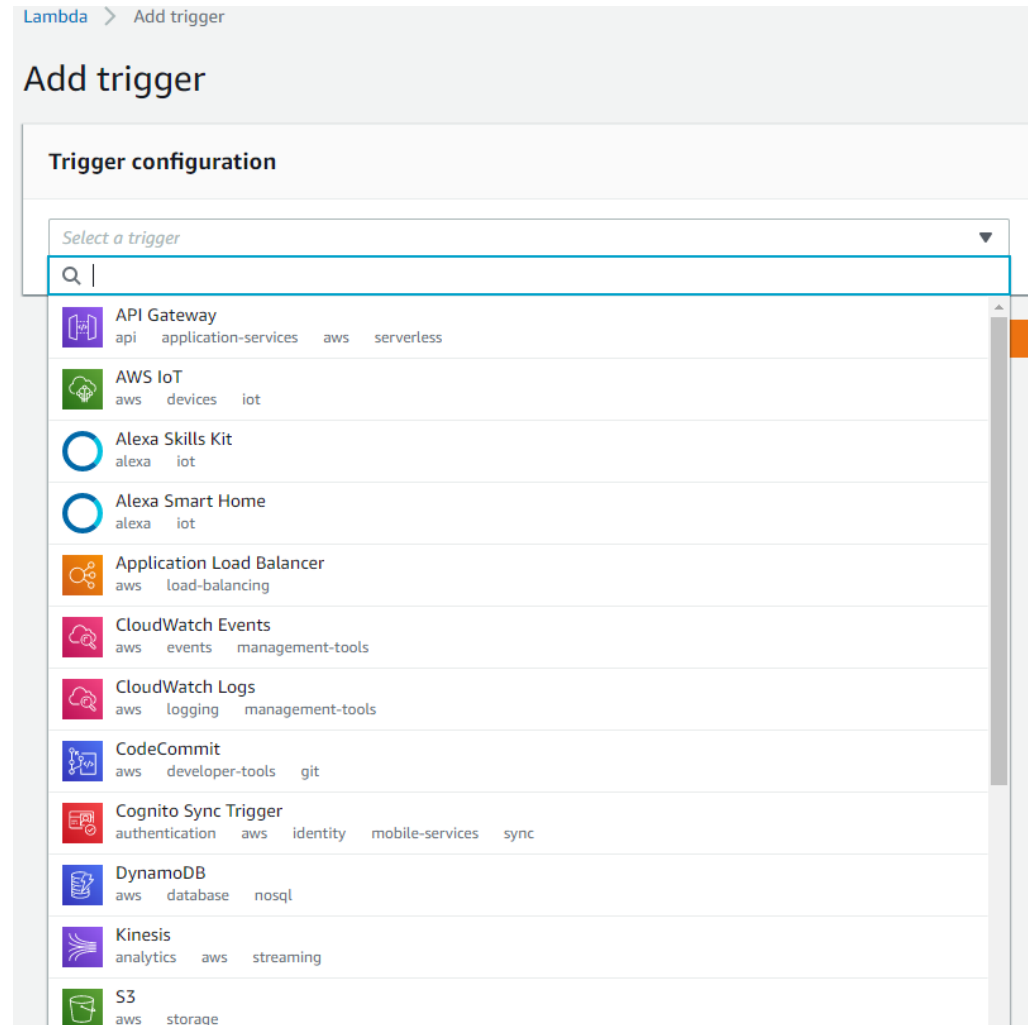
# Creating a Lambda Function

- The interface has changed quite a lot over the years
- First specify function name and execution environment
  - e.g. Node, Ruby, Python etc
- **Invokes Cloud9 Editor**
  - Online IDE
  - Acquired by AWS in 2016
- **Allows function to be input and simple tests to be performed**



# Next step is to add Event Triggers

- There are a large selection of things happening within AWS environment that can trigger the calling of a Lambda Function
- If programming a web REST API, then API gateway will field the HTTP requests and convert them into events before calling the Lambda function



# Using API Gateway as Event Trigger

<https://www.youtube.com/watch?v=DSrg7hG-jV4&list=PLzvRQMj9HDiSQMe68cti8cupl0mzLk1Gc&index=3>

[5 min]

# Result of invoking API to get event data

```
{
  "resource": "/DOMfirstfunc",
  "path": "/DOMfirstfunc",
  "httpMethod": "GET",
  "headers": {
    "accept": "text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3",
    "accept-encoding": "gzip, deflate, br",
    "accept-language": "en-US,en;q=0.9",
    "Host": "4y94vcgivg.execute-api.eu-west-1.amazonaws.com",
    "sec-fetch-mode": "navigate",
    "sec-fetch-site": "none",
    "upgrade-insecure-requests": "1",
    "User-Agent": "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/77.0.3865.120 Safari/537.36",
    "X-Amzn-Trace-Id": "Root=1-5dbf2bc5-d40b3965baba4de43211824d",
    "X-Forwarded-For": "37.228.252.43",
    "X-Forwarded-Port": "443",
    "X-Forwarded-Proto": "https",
    "multiValueHeaders": {
      "accept": ["text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3"],
      "accept-encoding": ["gzip, deflate, br"],
      "accept-language": ["en-US,en;q=0.9"],
      "Host": ["4y94vcgivg.execute-api.eu-west-1.amazonaws.com"],
      "sec-fetch-mode": ["navigate"],
      "sec-fetch-site": ["none"],
      "upgrade-insecure-requests": ["1"],
      "User-Agent": ["Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/77.0.3865.120 Safari/537.36"],
      "X-Amzn-Trace-Id": ["Root=1-5dbf2bc5-d40b3965baba4de43211824d"],
      "X-Forwarded-For": ["37.228.252.43"],
      "X-Forwarded-Port": ["443"],
      "X-Forwarded-Proto": ["https"]
    },
    "queryStringParameters": null,
    "multiValueQueryStringParameters": null,
    "pathParameters": null,
    "stageVariables": null,
    "requestContext": {
      "resourceId": "i2tmc4",
      "resourcePath": "/DOMfirstfunc",
      "httpMethod": "GET",
      "extendedRequestId": "CmPG7EDZDoEF8KA=",
      "requestTime": "03/Nov/2019:19:34:29 +0000",
      "path": "/prod/DOMfirstfunc",
      "accountId": "989357430853",
      "protocol": "HTTP/1.1",
      "stage": "prod",
      "domainPrefix": "4y94vcgivg",
      "requestTimeEpoch": 1572809669872,
      "requestId": "a904c3ec-3cce-4d92-ad9f-6018bf6ede9a",
      "identity": {
        "cognitoIdentityPoolId": null,
        "accountId": null,
        "cognitoIdentityId": null,
        "caller": null,
        "sourceIp": "37.228.252.43",
        "principalOrgId": null,
        "accessKey": null,
        "cognitoAuthenticationType": null,
        "cognitoAuthenticationProvider": null,
        "userArn": null,
        "userAgent": "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/77.0.3865.120 Safari/537.36",
        "user": null
      },
      "domainName": "4y94vcgivg.execute-api.eu-west-1.amazonaws.com",
      "apiId": "4y94vcgivg",
      "body": null,
      "isBase64Encoded": false
    }
  }
}
```



# Lambda Pricing

- **Charges are based on the number of invocations and the intensity of each invocation**
- **Invocations:**
  - First 1 million invocations/month are free
  - \$0.20 PER 1M Requests thereafter: \$0.0000002 per request.
- **Intensity**
  - measured in GB-seconds
    - product of memory allocated and seconds of CPU time
  - First 400,000 GB-seconds per month is free
  - \$0.0000166667 for every GB-SECOND used thereafter
  - Minimum memory of 128MB per instance
- **S3 or DynamoDB charges are separate**

# Pricing Example

If your Lambda@Edge function executed 10 million times in one month, and it ran for 50ms each time, your charges would be calculated as follows:

## Monthly compute charges

The monthly compute price is \$0.00000625125 per 128MB-second

Total compute (seconds) =  $10M * (0.05\text{sec}) = 500,000$  seconds

Monthly compute charges =  $500,000 * \$0.00000625125 = \$3.13$

## Monthly request charges

The monthly request price is \$0.60 per 1 million requests.

Monthly request charges =  $10M * \$0.6/M = \$6.00$

## Total monthly charges

Total charges = Compute charges + Request charges =  $\$3.13 + \$6.00 = \$9.13$   
per month

# Serverless Lambda Function-as-a-Service FaaS Summary

- **Intention is to package up useful work as a 'Function'**
- **Cloud environment facilitates connecting events to the function (web, IoT, Messaging)**
- **User does not have to worry about compute resources or scaling**
- **Pricing is based on function invocations i.e. useful work done**

# Serverless: Function-as-a-Service FaaS Differences from Platforms-as-a-Service

- **Both hide "servers" from developers**
- **PaaS typically always has at least one server process running that receives external requests**
- **scaled by booting up more server processes**
  - charged for
- **FaaS no server process constantly being run**
  - ~pay for function execution time
  - no process idle time

# Alternative Approach, manage the hardware like software

- **Infrastructure as code (IaC)**
  - machine-readable definition files
  - managing and provisioning computer data center resources
    - AWS CloudFormation
      - JSON or YAML
    - Terraform
      - HashiCorp Configuration Language (HCL), or JSON
      - August 2023 moved from Open source to source available
        - » Business Source License